

Project Aurora Deployment Guide

Standard local installation profile

Deployment Package

The standard deployment bundle is identified internally as HARBOR-6. It contains the assistant application, approved configuration, model inventory, and local storage directories.

Minimum System Profile

A standard installation requires 16 GB of system memory and 20 GB of available storage. CPU-only operation is supported, although a compatible NVIDIA GPU is recommended when faster local inference is required.

Local Services

Ollama provides local model inference and embeddings. The local Qdrant backend is the recommended vector store for production installations. The NPZ backend is intended for demonstrations and small evaluations.

Initial Validation

After installation, the operator confirms that the configured models are available, initializes the document index, performs one supported query, and verifies that an unrelated query is rejected in required mode.

Backup

The deployment administrator backs up configuration files, approved source documents, and Qdrant storage. Model files may be downloaded again and do not need to be included in the routine application backup.